Systematic Architecture and Parameterization of Quantitative Classifier Systems: A Rigorous Evaluation of SRS Open Questions

Executive Overview

The formulation of a highly reliable quantitative classifier system necessitates the precise calibration of numerous interconnected computational modules, spanning algorithmic classification, composite scoring mechanisms, implied volatility (IV) dislocation models, and execution decision engines. The rigorous parameterization of the System Requirements Specification (SRS) requires resolving critical dependencies surrounding long-horizon data stationarity, error tolerances in cumulative distribution functions, data fusion latency across multi-sensor arrays, and the complex mathematical modeling of option Greeks—specifically theta decay—in short-duration convex structures.[1]

This comprehensive research report provides the theoretical and empirical foundations required to systematically update the SRS, Architecture Decision Records (ADRs), and internal registry configurations. The ensuing analysis addresses specific mathematical and architectural questions spanning the Classifier Formula (CLS-001), Composite Scoring (CLS-002), IV Dislocation (CLS-006), the Decision Engine (DEC-001), and the overarching necessity of formal Regime Detection models.[1] By synthesizing advanced multi-sensor data fusion principles, stochastic volatility modeling, Bayesian inference for missing data, and high-frequency trading microstructure, optimal configurations for algorithmic deployment are established. The objective is to transition the architecture from a rigid, historically curve-fitted framework into an adaptive, regime-aware deployment engine capable of surviving and exploiting persistent market volatility.

Classifier Formula Architecture (CLS-001): Distributional Stationarity and Encoding

The integrity of any quantitative trading signal relies fundamentally on the statistical robustness of its underlying inputs. The structural design of the Classifier Formula dictates how raw market data is transformed into standardized, actionable scores. A central challenge in this domain is ensuring that the statistical methods applied to historical data remain valid in current and future market states.

Long-Horizon ECDF Stationarity and Volatility Regimes

A foundational question regarding the Classifier Formula concerns the assumption of

stationarity in the distribution of the daily absolute deviation from the median for the Cboe Volatility Index (VIX) over a five-year horizon (N_L).[1] The VIX, commonly referred to as the "fear gauge," measures the 30-day expected volatility of the S&P 500 Index derived from a strip of SPX options expiring between 23 and 37 days from the current trading day.[2] Since its inception, the VIX has historically averaged 19.38% with a standard deviation of 7.9%.[2] While the VIX exhibits well-documented mean-reverting properties over the long term, its absolute deviation is strictly non-stationary across any fixed five-year horizon.[5]

Financial markets transition through distinct, persistent states of calm and turbulence, creating structural breaks in volatility distributions.[6] A five-year lookback inherently encompasses multiple, distinct market regimes. For instance, the historically suppressed volatility environment of 2017 contrasts sharply with the extreme leptokurtic shocks experienced during the 2020 pandemic crisis or the 2008 financial crisis.[1] During stable regimes, persistently low volatility often leads market participants to take on excessive risk, ultimately seeding the conditions for violent expansions—a systemic risk phenomenon defined in the literature as the "volatility paradox".[8]

Consequently, utilizing a static Empirical Cumulative Distribution Function (ECDF) over a fixed five-year period essentially merges distinct statistical distributions into a single, flawed aggregate. This aggregation leads to a mischaracterization of the current market state, typically resulting in an overestimation of baseline risk during prolonged calm periods, and a severe underestimation of tail risk during systemic shocks.[9] Because volatility clusters and exhibits long memory, the absolute deviation from the median shifts dynamically, rendering the stationarity assumption mathematically invalid over a multi-year horizon.[6]

Minimum ECDF Window and Standard Error Calibration

The ECDF is a non-parametric estimator that evaluates the probability $P(X \leq x)$ directly from sample data without assuming an underlying parametric distribution (such as a Gaussian or log-normal distribution).[11] The standard error (SE) of the ECDF percentile estimate is determined by the binomial variance of the underlying proportion. For an estimated cumulative probability $p = F_n(x)$, the standard error is given by the formula:

$$SE = \sqrt{\frac{p(1-p)}{n}}$$

A critical requirement of the system architecture is to determine the minimum long-horizon window size (N_L) required to suppress the ECDF percentile standard error strictly below 0.03.[1] To guarantee that this condition is met across the entire distribution, the maximum variance scenario must be evaluated. The variance of a binomial proportion is mathematically maximized

at the median, where $p = 0.5$.[12]

Solving for the minimum sample size N_L:

$$0.03 > \sqrt{\frac{0.5 \times (1 - 0.5)}{N_L}}$$

$$0.03 > \frac{0.5}{\sqrt{N_L}}$$

$$\sqrt{N_L} > \frac{0.5}{0.03}$$

$$\sqrt{N_L} > 16.667$$

$$N_L > 277.78$$

Therefore, the absolute minimum sample size required to achieve a standard error below 0.03 at the median is 278 observations.[12] In the context of daily trading data, where one calendar year comprises approximately 252 trading days, an N_L of 278 represents roughly 1.1 years of daily closing data. A five-year daily lookback provides approximately $N = 1260$ observations. At this sample size, the maximum standard error at the median is compressed to approximately 0.014, satisfying the 0.03 architectural requirement effortlessly and providing highly precise percentile estimates for daily data like the VIX.[12]

However, a critical mathematical degradation occurs when this identical framework is applied to lower-frequency macroeconomic indicators.[1] For monthly macroeconomic indicators, such as the year-over-year Consumer Price Index (CPI_YOY), a five-year window yields only $N = 60$ observations. For weekly indicators, such as Initial Jobless Claims, a three-year window yields $N = 156$ observations.[1]

Table 1 illustrates the maximum standard errors for varying lookback periods and sampling frequencies, demonstrating the limitations of the ECDF approach for sparse datasets.

Indicator Frequency	Lookback Period	Sample Size (N)	Max Standard Error (p=0.5)	Tail Standard Error (p=0.95)	Meets SE<0.03 Threshold
Daily (e.g., VIX)	1.1 Years	278	0.0300	0.0130	Yes
Daily (e.g., VIX)	5 Years	1260	0.0141	0.0061	Yes
Weekly (e.g., Claims)	3 Years	156	0.0400	0.0174	No (fails at median)
Monthly (e.g., CPI)	5 Years	60	0.0645	0.0281	No (fails at median)

The calculations reveal that sample sizes of $N = 60$ and $N = 156$ fundamentally fail to meet the $SE < 0.03$ threshold at the median.[12] The implications for the classifier system are profound: relying on a non-parametric ECDF for monthly macroeconomic data over a five-year horizon introduces a margin of error exceeding 6.45%. To rectify this structural flaw, the architecture must implement one of two solutions. The first is to extend the lookback horizon for low-frequency data to at least 23.1 years ($N = 278$ months) to accumulate sufficient samples. However, analyzing economic data over a quarter-century introduces severe macroeconomic regime shifts, rendering the historical data largely irrelevant to the current state. The superior alternative is to abandon the non-parametric ECDF for sparse data in favor of fitting a parametric distribution (such as a Gaussian or log-normal distribution). By estimating only the mean and variance from the 60 observations, the system can derive cumulative probabilities parametrically, significantly reducing the standard error of the percentile estimates, provided the underlying data does not severely violate the distributional assumptions.[11]

Score Schema: Absolute versus Signed Values

The structural encoding of signal severity determines how downstream modules process the data. An open question exists regarding whether the severity score should remain an absolute value or become a signed value [-1, +1] to preserve directionality.[1] Utilizing an absolute value for the score effectively destroys critical directional information, mathematically mapping both massive standard deviation expansions (a bullish volatility shock) and massive standard deviation compressions (a bearish volatility implosion) to the same positive vector.

Transitioning the score schema to a signed value [-1, +1] is imperative for advanced systematic operations.[1] Modern quantitative alpha models depend heavily on directionality to generate specific predictions regarding future price actions and to manage risk efficiently.[19] A signed schema allows the decision engine to natively differentiate between long-volatility and short-volatility conditions, seamlessly switching between mean-reversion and trend-following algorithms based on the vector of the signal.[21]

The downstream implications of adopting a signed schema are vast and universally beneficial. In composite scoring, signed values enable opposite signals from non-correlated assets to cleanly cancel each other out, accurately reflecting a state of market neutrality or directional conflict. Conversely, an absolute value schema would falsely aggregate opposing forces into a high composite score, erroneously suggesting high conviction when the true underlying state is high noise. Furthermore, in IV dislocation calculations, signed scores preserve the specific vector of the dislocation, explicitly defining whether implied volatility is trading at an excessive premium or a discount relative to realized volatility, which dictates whether the optimal strategy is to sell or buy premium.[22]

Regime Detection Architecture: Mathematical Necessity over Implicit Proxies

The question of whether a formal regime detection model is strictly necessary, or if regime context can be implicitly handled by a long-horizon ECDF, is central to the system's structural integrity and its ability to adapt to shifting macroeconomic conditions.[1]

The Insufficiency of Implicit ECDF Context and Rolling Percentiles

Implicitly handling regime context via a long-horizon ECDF relies on the dangerous assumption that extreme events will naturally and permanently fall into the upper percentiles of the historic distribution. However, this method exhibits significant temporal latency and memory-decay flaws.[25] When the market transitions abruptly from a low-volatility environment to a high-volatility crash, the initial stages of the crash are correctly identified as statistically extreme by the historical ECDF. Yet, as the high-volatility regime persists over weeks or months, the rolling ECDF slowly absorbs the new extreme data, expanding its variance. This gradually dampens the severity scores of what are objectively highly volatile days because they have mathematically become the "new normal" within the shifting window. This structural lag causes momentum and trend-following strategies to misallocate capital during regime transitions, either

entering too late or failing to exit when the market character shifts.[6]

Furthermore, standard rolling volatility percentiles and basic K-Means clustering algorithms typically fail to capture higher-order structural transformations, often defaulting to binary "bullish" or "bearish" classifications that lack the nuance required for sophisticated derivatives trading.[27]

Hidden Markov Models (HMM) and Adaptive Hierarchical Structures

A formal regime detection model, particularly a Hidden Markov Model (HMM), is mathematically superior to simple rolling percentiles and is strictly required for the classifier architecture. HMMs do not assume fixed transition probabilities; rather, they assume the market exists in a discrete number of unobservable (latent) states (e.g., Equity Trend, Normal, Rate Shock, Volatility Shock).[27]

The market's observable variables—such as logarithmic returns, variance, bid-ask spreads, and macroeconomic indicators—are emitted by the latent state according to state-specific probability distributions.[10] By applying the Baum-Welch algorithm or utilizing Adaptive Hierarchical Hidden Markov Models (AH-HMM), the system continuously computes the probability of occupying each state and the likelihood of transitioning to another.[7] AH-HMMs introduce a meta-regime layer that reflects the broader macro-financial environment, allowing the transition matrix itself to adapt dynamically to structural changes arising from crises or policy shifts.[7]

This approach detects structural breaks almost immediately, analyzing the "muffled bits of music leaking through the wall" to infer the hidden state of the market.[27] Empirical research demonstrates that HMMs significantly outperform static strategies and simple rolling windows in both in-sample likelihood and out-of-sample Value-at-Risk coverage, identifying key turning points like the Global Financial Crisis and the COVID-19 shock with higher precision.[7]

Binding System Variables to Regime State

Because market behavior alters fundamentally across different states, specific parameters within the decision engine must be strictly regime-dependent.[29] A volatility breakout signal generated in a historically low-volatility regime possesses completely different predictive power, expected duration, and mathematical expectation values than the exact same signal generated in a high-volatility regime.[6]

Consequently, capital sizing, option holding periods, and execution thresholds cannot remain static.[6] HMMs provide the conditional probabilities necessary to dynamically adjust these parameters. For instance, momentum and trend-following strategies generally perform better in stable regimes, while mean-reversion and contrarian strategies dominate volatile environments.[6] The system must utilize the output of the HMM to dynamically adjust the SENSITIVITY_FACTOR in the volatility models and shift the allocation weightings in real-time based on the prevailing regime.[6]

Composite Scoring (CLS-002): Multi-Sensor Data

Fusion and Noise Suppression

To produce high-conviction deployment triggers, signals from multiple disparate indicators must be aggregated into a single composite score. This requires carefully balancing noise suppression techniques against the risk of missing early, explosive market movements.[1]

Conflict-Based Weighting and Corroboration Thresholds

In quantitative trading environments, false positives generate severe frictional costs through unnecessary execution fees and bid-ask spread crossing, while false negatives result in missed alpha generation during critical market dislocations.[30] Drawing from multi-sensor data fusion principles traditionally utilized in robotics, autonomous navigation, and signal processing, conflict-based weighting offers an optimal mathematical solution for signal corroboration.[32]

Requiring corroboration strictly within the same source category severely limits the system's ability to act on valid, idiosyncratic order flow or localized liquidity vacuums. Conversely, requiring corroboration from any category introduces the risk of conflating unrelated macroeconomic events with microstructural noise. The optimal architectural solution is a dual-pathway threshold-waived corroboration model.

If a single instrument's signal exceeds a predefined, extreme high-conviction threshold (T)—representing an anomaly so statistically significant (e.g., a 4σ deviation or a massive options sweep) that it mathematically eclipses background noise—corroboration requirements are immediately waived.[1] This permits the instantaneous capture of flash events and liquidity dislocations where waiting for a secondary signal would result in the repricing move having already occurred.[26] For signals below the high-conviction threshold T, adaptive noise fusion strategies must be applied. By utilizing covariance matrices to measure signal consensus and down-weighting isolated sensors that exhibit poor consensus (conflict measures), the system yields improved accuracy over naive averaging, effectively suppressing impulse noise and Gaussian corruption without discarding valuable data.[32]

Temporal Asymmetry in Corroboration Windows

The acceptable latency between an initial signal and a corroborating signal must be strictly bifurcated based on the underlying event typology.[1] A uniform latency window is incompatible with the diverse nature of financial data.

1. **Intraday Event-Driven Signals (Micro-Trend):** For order-book imbalances, sudden volatility spikes, or algorithmic flash crashes, the market digests information in milliseconds to microseconds.[35] In these ultra-low latency environments, the corroboration latency window must be tightly constrained (e.g., 100 to 300 milliseconds). If an options order book flashes a massive structural imbalance, but the underlying asset's pricing feed does not reflect a corresponding tick adjustment within 300 milliseconds, the signal is

highly likely to be structural noise, a localized liquidity void, or a spoofing attempt.[35]

2. **Macroeconomic Releases (Macro-Trend):** Conversely, macroeconomic data releases (e.g., CPI inflation data, Non-Farm Payrolls, Central Bank rate decisions) require significantly longer digestion periods. While high-frequency algorithms react instantaneously to the headline numbers, broader market repositioning, institutional capital flows, and the alignment of multi-timeframe directional biases take time to formulate.[25] Academic research indicates that inflation swap data and macro-driven market sentiment require an optimal observation window of approximately 30 to 60 minutes to fully incorporate the new information without being stopped out by whipsaw price action.[25] During macroeconomic shocks, prices often spike, reverse sharply as mean reversion takes over, and then establish a true trend.[26] Therefore, macro-corroboration windows should permit up to 60 minutes for secondary indicators and broader market breadth to align, effectively treating macro hours as special market conditions.[26]

Advanced Imputation Strategies for Source-Dropout Penalties

When an expected data source is unavailable, applying a default static scalar penalty of 0.7 to the composite score is an overly simplistic heuristic.[1] This figure is typically derived from standard availability engineering concepts, where system reliability is often estimated against expected uptime, but it lacks empirical foundation in financial econometrics.[41] In quantitative finance, missing data is rarely ignorable; it is usually Missing Not At Random (MNAR).[43]

If the CBOE volatility feed or a major pricing API goes down, it does not mean the market is statically 30% less reliable. Outages often correlate with immense market stress, extreme volume spikes, or systemic infrastructure failure. Instead of a flat 0.7 penalty, the penalty function should scale logarithmically based on the duration of the outage. A microsecond dropout is routine UDP packet loss; a multi-minute dropout is indicative of a systemic event requiring a severe degradation of the composite confidence score.[44]

Furthermore, rather than simply penalizing the composite score, the architecture should employ Bayesian updating and imputation to reconstruct the missing data.[45] By treating the missing source as a latent variable and utilizing the covariance of currently available, highly correlated sensors, the system can mathematically impute the missing value using prior distributions. This approach dynamically adjusts the uncertainty (and thus, the weight of the penalty) based on the variance of the imputation, blending what is already known with what the active sensors are currently indicating.[43]

IV Dislocation Analytics (CLS-006): Data Acquisition and Sensitivity Calibration

Implied Volatility (IV) dislocation models require low-latency, highly accurate data to measure discrepancies between the market's expectation of volatility (derived from option prices) and the underlying asset's historical or expected realized volatility.[49] The calculation of these

discrepancies is highly sensitive to the inputs. Relying on stale daily closes from free public sources like the Federal Reserve Economic Data (FRED) introduces unacceptable latency that destroys algorithmic edge.[1]

High-Frequency Market Data Engineering and Staleness Tolerance

Access to intraday volatility data—specifically the VIX and the Cboe Crude Oil ETF Volatility Index (OVX)—is a critical path requirement for the architecture.[51] While standard retail Broker APIs offer general market data, they are frequently subject to rate throttling, arbitrary packet delays, and internal data aggregation that obscures true tick-by-tick liquidity.

For professional-grade applications, the choice typically resides between direct exchange connections and specialized tier-1 aggregators. Direct CBOE APIs and co-located exchange feeds operate in the microsecond (μs) to nanosecond (ns) regime, which is essential for pure high-frequency trading (HFT) and market making.[35] However, these direct feeds require massive infrastructure investments. Alternatively, tier-1 aggregators like Twelve Data provide a WebSocket API (wss://ws.twelvedata.com) for real-time streaming. This protocol circumvents the HTTP overhead of standard REST requests, delivering data with low millisecond latency.[52]

The acceptable staleness threshold for reliable dislocation calculations depends entirely on the option pricing model being utilized. Computing IV requires numerical methods, particularly the Newton-Raphson iterative procedure, because the Black-Scholes pricing formula is nonlinear in volatility and lacks a closed-form solution.[50] The Newton-Raphson method leverages Vega—the sensitivity of the option price to volatility—for rapid estimation.[50] If the underlying asset price data is stale by even a few seconds while the option premium is updating, the calculated IV will spike erroneously, triggering false dislocation signals. For strategies holding positions longer than a few minutes, a WebSocket feed ensuring the staleness threshold remains strictly under 100 milliseconds is sufficient, ensuring the Greeks calculated by the pricing model accurately reflect the immediate state of the order book.[35]

VIX and OVX Regime Mapping and Structural Variances

Properly mapping volatility regimes requires defining the specific breakpoints where market microstructure fundamentals shift.[1] Historical analysis of the VIX, which measures standard deviation expectations for the S&P 500, suggests four distinct behavioral regimes:

- **Low Volatility (Equity Trend):** 0 to 15. Characterized by high investor confidence, low option premiums, sustained bull markets, and slow, steady capital appreciation.[28]
- **Medium Volatility (Normal):** 15 to 25. The historical mean of 19.38% resides here. Represents a functioning, two-way market with standard price discovery and moderate fluctuations.[2]
- **High Volatility (Rate/Macro Shock):** 25 to 30 (or up to 35). Escalating turbulence, widening bid-ask spreads, and sharp mean-reverting price action. In this regime, VIX calls become highly active as institutions hedge portfolios.[28]

- **Extreme Volatility (Volatility Shock):** 30+. Driven by fear, panic liquidations, and massive tail-risk hedging. Represents severe market dislocations and the breakdown of normal asset correlations.[28]

Conversely, the OVX, which measures the 30-day expected volatility of crude oil as priced by the United States Oil Fund (USO ETF), inherently trades at a significantly higher baseline.[51] While equities trend upward over time with occasional sharp corrections, commodities are heavily influenced by supply chain disruptions, geopolitical conflicts, and inventory data, leading to a naturally higher state of constant variance.[60] The median value for OVX fluctuates around 32.8%, whereas the VIX median is substantially lower.[59]

Table 2 highlights the structural differences in regime breakpoints between broad equities (VIX) and commodities (OVX), demonstrating why a unified threshold across asset classes is mathematically unsound.

Regime Classification	VIX Breakpoints	OVX Breakpoints	Dominant Market Mechanics
Low Volatility	0 - 15	< 30	Directional trending, suppressed option premiums, high confidence.[57]
Medium / Normal	15 - 25	30 - 35	Baseline mean reversion, balanced order books, normal liquidity.[57]
High Volatility	25 - 30	35 - 50	Geopolitical tension, supply chain noise, rapid gap-ups, spread widening.[57]
Extreme / Shock	> 30	> 50 (up to 100+)	Panic liquidations, structural supply outages (e.g., pandemic, war), margin calls.[57]

Volatility Model Calibration and Dynamic Sensitivity Factors

The SENSITIVITY_FACTOR parameter dictates the responsiveness of the chosen implied volatility models (e.g., Stochastic Volatility Inspired (SVI), Heston, or Kennedy models) to

changes in the underlying asset.[1] It operates similarly to the Greek "Vega," tracking how option prices theoretically respond to a shift in implied volatility.[50]

Because volatility exhibits pronounced convexity—tending to rise much more dramatically during large equity losses than large equity gains—the SENSITIVITY_FACTOR cannot be a static scalar across all regimes.[63] In a low VIX regime (<15), the mathematical probability of a sudden volatility expansion is vastly higher than further compression, requiring an asymmetric upward sensitivity factor to accurately price tail risk.[29] Furthermore, stochastic volatility models with regime switching (such as an extended Heston model) demonstrate that failure to adapt the sensitivity parameter causes severe mispricing in the tails of the distribution.[65]

During high-volatility regimes (>30), the sensitivity factor must actually dampen to prevent the algorithmic engine from drastically over-hedging against transient noise spikes.[29] The calibration requires complex numerical approximations, extracting discrete parameter sensitivity through a central finite-difference approach.[63] This dynamic adjustment ensures that the algorithm does not chase noise during chaotic periods, while remaining hypersensitive to early warning signs of dislocation during periods of calm.

Decision Engine Kinetics (DEC-001): Threshold Optimization and Convex Trade Management

The Decision Engine acts as the final arbiter and terminal node of the quantitative architecture. It takes the normalized, regime-adjusted composite scores from the classifier and executes deterministic capital deployment.[1]

Objective Functions and Tiered Deployment Thresholds

Defining the optimal composite score threshold for deployment relies on balancing the inherent cost asymmetry between false positives (executing a trade that decays to zero) and false negatives (missing a lucrative tail-risk event).[1]

Implementing a binary threshold leads to brittle execution logic. A tiered execution protocol is mathematically and operationally superior. A lower composite score should trigger a preliminary "Watch" state. This state activates heightened polling frequencies from data sensors, warms up execution routes, and begins preliminary risk calculations. Only upon breaching a secondary "Deploy" threshold is capital actually allocated.[1]

A critical vulnerability in defining this ultimate deploy threshold is over-fitting to historical data—a phenomenon where machine learning systems learn spurious patterns in historical prices and fail spectacularly out of sample.[66] Strategies optimized purely on maximal historical return often suffer from data snooping bias.[66] To counter this, the deploy threshold must be calibrated using a composite objective function, such as the GT-Score. This function integrates performance, statistical significance, consistency, and downside risk to force the engine toward generalized robustness rather than curve-fitting the noise, substantially improving the generalization ratio in

walk-forward validations.[66]

The Mechanics of Theta Decay and Vega Crush in Short-Duration Holdings

When deploying non-directional, convex options strategies designed to profit from extreme volatility expansions—such as straddles (buying an at-the-money call and put) or strangles (buying an out-of-the-money call and put)—the overriding mathematical adversary is time decay, denoted by the Greek letter Theta (Θ).[54]

Theta measures the expected daily decrease in an option's extrinsic value, representing the inescapable cost of holding the position.[54] Crucially, Theta decay is non-linear; it accelerates aggressively as the option approaches expiration, creating a distinct "hockey stick" curve on a graphical plot.[54] The structural question for the SRS asks at what specific point daily theta decay completely negates the mathematical edge of holding a convex position for 6 to 72 hours.[1]

1. **The Strangle Dynamics:** A long strangle requires a massive directional move to surpass both the initial implied volatility premium paid and the large gap to the out-of-the-money strike price.[67] In a historically low implied volatility environment, they are exceptionally cheap to enter. However, if held for 72 hours with zero underlying price action, the accelerated theta on near-term options will ravage the premium, rendering the position practically worthless before expiration.[67]
2. **The Straddle Dynamics:** Straddles are significantly more expensive because both options are purchased at-the-money, meaning they carry the highest possible extrinsic value.[70] While they require a smaller absolute price move to break even compared to a strangle, their high extrinsic value subjects them to the highest absolute Theta exposure.[72]

However, for a 6 to 72-hour holding period—particularly leading into an anticipated catalyst such as an earnings report or a CPI print—the mathematical edge is rarely destroyed purely by Theta. The true destroyer of edge in short-duration convex holds is **Vega (Volatility Crush)**.[73]

If a straddle is purchased 24 hours prior to a known binary event, implied volatility is bid up by market participants to exorbitant levels.[73] Immediately following the event, the uncertainty collapses, causing IV to plummet. Even if the underlying asset price moves directionally in the trader's favor, the drop in IV (measured by Vega) can instantly obliterate the value of both legs of the straddle.[73] Positive Gamma (the rate of change of Delta) must overwhelm the negative impact of Vega for the trade to be profitable.[73]

Table 3 models the theoretical interaction between the options Greeks during a short-term catalyst event, demonstrating why hold periods must be strictly managed.

Market	Underlying	Implied	Theta	Vega	Net Option

Phase	Asset Move	Volatility (IV) State	Impact	Impact	Value
Pre-Event (T-24h)	Minimal	Expanding	Moderate drain	Highly Positive	Appreciating
Event Horizon	None	Peaked	Severe drain	Neutral	Decaying
Post-Event (Miss)	< Implied Move	Collapsing (Crush)	Severe drain	Highly Negative	Massive Loss [73]
Post-Event (Hit)	> Implied Move	Collapsing (Crush)	Severe drain	Negative	Profitable (Gamma overtakes Vega) [73]

The mathematical conclusion for the Decision Engine is strict: for short-duration hold periods (6 to 72 hours), positions must generally be exited *prior* to the binary event if the original trade thesis relied primarily on IV expansion leading up to the catalyst.[74] If the intention is to hold through the event, the directional movement threshold required for a post-event hold must be dynamically calculated to significantly exceed the modeled post-event volatility crush.[73] If the expected Gamma from the asset move cannot demonstrably overwhelm the anticipated drop in Vega, the decision engine must reject the deployment.

Portfolio Heat, Global Caps, and Opposing Signal Resolution

Once a position is opened, subsequent signals may fire from the classifier. If a new signal matches the current directional bias, the decision to add to the position must be mathematically weighed against portfolio heat and overall risk exposure.[1] If the new signal opposes the current open position, the engine faces a conflict.[1] In a signed score schema [-1, +1], an opposing signal acts as a structural stop-loss or profit-taking mechanism, prompting a reduction in delta exposure rather than opening a completely new, friction-heavy opposing position which would incur unnecessary bid-ask spread costs.[20]

Finally, a global position size cap is mathematically non-negotiable.[1] Because fully automated quantitative systems can occasionally enter positive feedback loops during unprecedented market regimes (such as the recursive selling algorithms involved in the 1987 Black Monday crash or the 2010 Flash Crash), a hardcoded global risk limit prevents catastrophic capital ruin. Allocation constraints must enforce that the total portfolio at-risk capital never exceeds the maximum drawdown tolerance calculated by the Monte Carlo simulations of the system's worst historical underperformance across thousands of random seed scenarios.[66]

Synthesis and Architectural Directives

The exhaustive investigation into the open parameters of the classifier system yields several definitive structural resolutions that must be codified into the System Requirements Specification (SRS):

1. **Classifier Configurations:** A minimum long-horizon window size (N_L) of 278 observations is mathematically required to secure an ECDF standard error below the 0.03 threshold. Due to this constraint, weekly and monthly macroeconomic indicators cannot be processed via a non-parametric 5-year ECDF and must be transitioned to parametric fitting to avoid unacceptable margins of error. Furthermore, the scoring schema must immediately migrate from absolute values to a signed [-1, +1] topology to preserve necessary directional vectors and covariance data for downstream alpha models.[1]
2. **Signal Fusion and Corroboration:** The system must abandon static category corroboration in favor of adaptive multi-sensor data fusion. Signals exceeding a predefined standard deviation anomaly threshold (T) will bypass corroboration entirely to capture low-latency liquidity dislocations. Routine signals require corroboration within 300 milliseconds for intraday microstructure events, and up to 60 minutes for macroeconomic trend shifts. The static 0.7 source dropout penalty must be replaced with dynamic Bayesian imputation.[1]
3. **Regime Detection Models:** Implementing an Adaptive Hierarchical Hidden Markov Model (AH-HMM) is strictly required. The baseline assumption that a rolling ECDF can implicitly handle regime changes is empirically false due to unacceptable temporal latency and memory decay during structural breaks. All downstream thresholds, sizing metrics, and sensitivity parameters must conditionally bind to the current latent HMM state.[6]
4. **Implied Volatility Infrastructure:** Data pipelines should utilize tier-1 WebSocket APIs to ensure low-millisecond latency for IV dislocation models, avoiding the pitfalls of stale daily pricing. Breakpoints for regime analysis must be separated by asset class, specifically noting the heavily elevated baseline volatility of commodities (OVX, median ~32.8%) versus broad equities (VIX, median ~19.38%).[51]
5. **Execution Protocols:** The Decision Engine must employ a tiered Watch/Deploy threshold optimized via multi-variable objective functions to avoid historical data snooping. For 6 to 72-hour hold periods in convex options structures, the engine must explicitly calculate the theoretical post-event volatility crush (Vega risk), as Theta alone is secondary to the catastrophic impact of collapsing implied volatility following a known catalyst. Hardcoded global risk caps must be implemented to prevent recursive algorithmic feedback loops.[66]

By formally integrating these mathematical limits, latency requirements, and architectural transitions into the SRS, the algorithmic framework will achieve the necessary robustness to operate successfully across shifting financial paradigms.

Works cited

1. open-questions-2026-04-27.md
2. The CBOE VIX Fear Index - Current Market Valuation, accessed April 27, 2026, https://www.currentmarketvaluation.com/models/vix-fear-index.php
3. What the VIX and VIX1D Indices Attempt to Measure and How They Differ | Cboe, accessed April 27, 2026, https://www.cboe.com/insights/posts/what-the-vix-and-vix-1-d-indices-attempt-to-measure-and-how-they-differ/
4. Cboe Volatility Index Mathematics Methodology, accessed April 27, 2026, https://cdn.cboe.com/resources/indices/Cboe_Volatility_Index_Mathematics_Methodology.pdf
5. A Practitioner's Guide to Reading VIX - S&P Global, accessed April 27, 2026, https://www.spglobal.com/spdji/en/education-a-practitioners-guide-to-reading-vix.pdf
6. Volatility and Market Regimes: How Changing Risk Shapes Market Behavior (with Python Examples) | by Trading Dude | Medium, accessed April 27, 2026, https://medium.com/@trading.dude/volatility-and-market-regimes-how-changing-risk-shapes-market-behavior-with-python-examples-190de97917d8
7. Adaptive Hierarchical Hidden Markov Models for Structural Market Change - MDPI, accessed April 27, 2026, https://www.mdpi.com/1911-8074/19/1/15
8. Regime shifts in stock market volatility: a historical perspective on the US market, accessed April 27, 2026, https://www.ecb.europa.eu/press/financial-stability-publications/fsr/focus/2018/pdf/ecb~bcaaae16c3.fsrbox201805_03.pdf
9. Hidden Markov Model Market Regimes: How HMM Detects Market Regimes in Trading Strategies - QuantifiedStrategies.com, accessed April 27, 2026, https://www.quantifiedstrategies.com/hidden-markov-model-market-regimes-how-hmm-detects-market-regimes-in-trading-strategies/
10. Regime Detection and Risk Allocation Using Hidden Markov Models – BSIC, accessed April 27, 2026, https://bsic.it/regime-detection-and-risk-allocation-using-hidden-markov-models/
11. Empirical Cumulative Distribution Function - with solution, calculates the ECDF of sample data, draws chart. - Statistics Kingdom, accessed April 27, 2026, https://www.statskingdom.com/ecdf.html
12. Standard error - Wikipedia, accessed April 27, 2026, https://en.wikipedia.org/wiki/Standard_error
13. What Is Standard Error? | How to Calculate (Guide with Examples) - Scribbr, accessed April 27, 2026, https://www.scribbr.com/statistics/standard-error/
14. Determining sample size based on confidence and margin of error | AP Statistics | Khan Academy - YouTube, accessed April 27, 2026, https://www.youtube.com/watch?v=VyFs7fsWE6w
15. 7.5 Calculating the Sample Size for a Confidence Interval - eCampusOntario Pressbooks, accessed April 27, 2026, https://ecampusontario.pressbooks.pub/introstats/chapter/7-5-calculating-the-sam

ple-size-for-a-confidence-interval/
16. Summary of Minimum Sample Size Recommendations for SEM, accessed April 27, 2026, https://web.pdx.edu/~newsomj/semclass/ho_sample%20size.pdf
17. How to calculate minimum sample size when using ecdf? - Cross Validated, accessed April 27, 2026, https://stats.stackexchange.com/questions/365992/how-to-calculate-minimum-sample-size-when-using-ecdf
18. Sample size to detect decrease in standard deviation - Cross Validated, accessed April 27, 2026, https://stats.stackexchange.com/questions/370984/sample-size-to-detect-decrease-in-standard-deviation
19. What are Alpha Signals & How are They Used in Trading?, accessed April 27, 2026, https://www.tejwin.com/en/insight/alpha-signal/
20. Decoding 101 Financial Alpha Strategies: A Parsing Approach | by Jakub Polec | Medium, accessed April 27, 2026, https://medium.com/@jpolec_72972/decoding-101-financial-alpha-strategies-a-parsing-approach-02b997f7c36d
21. Volatility Index Trading: A Guide to VIX Market Success, accessed April 27, 2026, https://tradewiththepros.com/volatility-index-trading/
22. Alpha - QuantConnect.com, accessed April 27, 2026, https://www.quantconnect.com/docs/v2/writing-algorithms/algorithm-framework/alpha/key-concepts
23. AlphaEval: A Comprehensive and Efficient Evaluation Framework for Formula Alpha Mining - arXiv, accessed April 27, 2026, https://arxiv.org/pdf/2508.13174
24. Harvesting Volatility Risk Premium - Imperial College London, accessed April 27, 2026, https://www.imperial.ac.uk/media/imperial-college/faculty-of-natural-sciences/department-of-mathematics/math-finance/Shibo_Lu_01210524.pdf
25. Intraday and Post-Market Investor Sentiment for Stock Price Prediction: A Deep Learning Framework with Explainability and Quantitative Trading Strategy - MDPI, accessed April 27, 2026, https://www.mdpi.com/2079-8954/13/5/390
26. When Macro Hits: Why Quants Struggle and How to Adapt | by Oraichain Labs - Medium, accessed April 27, 2026, https://oraichain.medium.com/when-macro-hits-why-quants-struggle-and-how-to-adapt-811a776bcce4
27. Detecting Market Regimes: Hidden Markov Model | by David Borst - Medium, accessed April 27, 2026, https://datadave1.medium.com/detecting-market-regimes-hidden-markov-model-2462e819c72e
28. Built a simple CBOE vs VIX framework. Looking for feedback on methodology. - Reddit, accessed April 27, 2026, https://www.reddit.com/r/algotrading/comments/1s3ozj9/built_a_simple_cboe_vs_vix_framework_looking_for/
29. REGIME-DEPENDENT VOLATILITY SENSITIVITY - ResearchGate, accessed April 27, 2026,

https://www.researchgate.net/publication/404010679_REGIME-DEPENDENT_VOLATILITY_SENSITIVITY
30. Alpha Decay: what does it look like? And what does it mean for systematic traders?, accessed April 27, 2026, https://www.mavensecurities.com/alpha-decay-what-does-it-look-like-and-what-does-it-mean-for-systematic-traders/
31. Assessing the importance of the choice threshold in quantifying market risk under the POT approach (EVT) - PMC, accessed April 27, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC9818059/
32. Noise Fusion Strategy: Principles & Methods - Emergent Mind, accessed April 27, 2026, https://www.emergentmind.com/topics/noise-fusion-strategy
33. Multi-Sensor Data Fusion for Real-Time Multi-Object Tracking - MDPI, accessed April 27, 2026, https://www.mdpi.com/2227-9717/11/2/501
34. A Gentle Approach to Multi-Sensor Fusion Data Using Linear Kalman Filter - arXiv, accessed April 27, 2026, https://arxiv.org/abs/2407.13062
35. Latency Standards in Trading Systems - LuxAlgo, accessed April 27, 2026, https://www.luxalgo.com/blog/latency-standards-in-trading-systems/
36. Message Arrival Rates and Latency - Quantitative Trading, accessed April 27, 2026, https://markrbest.github.io/message-arrival-rates-and-latency/
37. Market Efficiency in Real Time: Evidence from Low Latency Activity around Earnings Announcements - MIT Sloan, accessed April 27, 2026, https://mitsloan.mit.edu/sites/default/files/inline-files/Chordia_Miao.pdf
38. The good, the bad, and latency: exploratory trading on Bybit and Binance - Taylor & Francis, accessed April 27, 2026, https://www.tandfonline.com/doi/full/10.1080/14697688.2025.2515933
39. Quantitative Trading Trends: Macro vs Micro Strategies | Quant Trading Platform - NURP, accessed April 27, 2026, https://nurp.com/algorithmic-trading-blog/macro-vs-micro-trends-in-quantitative-trading/
40. How Markets Process Macro News: The Importance of Investor Attention - Federal Reserve, accessed April 27, 2026, https://www.federalreserve.gov/econres/feds/files/2025022pap.pdf
41. System Reliability & Availability Calculations – BMC Software | Blogs, accessed April 27, 2026, https://www.bmc.com/blogs/system-reliability-availability-calculations/
42. Reliability optimization and system redesign under uncertain failure rates - Pure, accessed April 27, 2026, https://pure.tue.nl/ws/files/99259146/20180625_Ge.pdf
43. Bayesian Approaches for Missing Not at Random Outcome Data: The Role of Identifying Restrictions - PMC, accessed April 27, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC6936760/
44. Sequential Fusion Filter for State Estimation of Nonlinear Multi-Sensor Systems with Cross-Correlated Noise and Packet Dropout Compensation - PMC, accessed April 27, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC10224250/
45. Bayesian Missing Data Imputation — PyMC example gallery, accessed April 27, 2026,

https://www.pymc.io/projects/examples/en/latest/howto/Missing_Data_Imputation.html
46. Bayesian Statistics in Finance: A Trader's Guide to Smarter Decisions - Interactive Brokers, accessed April 27, 2026, https://www.interactivebrokers.com/campus/ibkr-quant-news/bayesian-statistics-in-finance-a-traders-guide-to-smarter-decisions/
47. Open Source Bayesian Models. 3. Composite Models for Prediction of Binned Responses | Journal of Chemical Information and Modeling - ACS Publications, accessed April 27, 2026, https://pubs.acs.org/doi/10.1021/acs.jcim.5b00555
48. Volume -1 Issue -1 | Bayesian Analysis - Project Euclid, accessed April 27, 2026, https://projecteuclid.org/journals/bayesian-analysis/volume--1/issue--1
49. Implied volatility - Wikipedia, accessed April 27, 2026, https://en.wikipedia.org/wiki/Implied_volatility
50. Implied Volatility: Formulation, Computation, and Robust Numerical Methods, accessed April 27, 2026, https://www.interactivebrokers.com/campus/ibkr-quant-news/implied-volatility-formulation-computation-and-robust-numerical-methods/
51. OVX Index Dashboard - Cboe Global Indices, accessed April 27, 2026, https://www.cboe.com/us/indices/dashboard/ovx/
52. API Documentation - Twelve Data, accessed April 27, 2026, https://twelvedata.com/docs
53. Twelve Data | Stock, Forex, and Crypto Market Data APIs, accessed April 27, 2026, https://twelvedata.com/
54. Theta Decay in Options Trading: 3 Strategies | Charles Schwab, accessed April 27, 2026, https://www.schwab.com/learn/story/theta-decay-options-trading
55. Understanding VIX or Volatility Index - TD Bank, accessed April 27, 2026, https://www.td.com/ca/en/investing/direct-investing/articles/understanding-vix
56. Volatility Index (VIX) Overview: What Is It, Formula, High vs. Low VIX - Gotrade, accessed April 27, 2026, https://www.heygotrade.com/en/blog/volatility-index-vix-overview/
57. What is the Volatility Index (VIX) — What you should know? - Trade Nation, accessed April 27, 2026, https://tradenation.com/articles/vix-volatility-index/
58. Volatility Index (VIX): Understanding Market Volatility and Types - ACY Securities, accessed April 27, 2026, https://acy.com/en/market-news/education/vix-market-volatility-index-and-types-111445/
59. Commodity Volatility Indices: OVX and GVZ | by Vito Turitto | HyperVolatility - Medium, accessed April 27, 2026, https://medium.com/hypervolatility/commodity-volatility-indices-ovx-and-gvz-e5ebc3dae501
60. Crude Awakening - Horizon Investments, accessed April 27, 2026, https://www.horizoninvestments.com/crude-awakening/
61. The SVI implied volatility model and its calibration - Diva-portal.org, accessed April 27, 2026, https://www.diva-portal.org/smash/get/diva2:744907/FULLTEXT02.pdf
62. Calibration and Option Pricing with stochastic volatility and double exponential

jumps - arXiv, accessed April 27, 2026, https://arxiv.org/html/2502.13824v1
63. An Analysis of Implied Volatility, Sensitivity, and Calibration of the Kennedy Model - MDPI, accessed April 27, 2026, https://www.mdpi.com/2227-7390/13/21/3396
64. VIX | S&P Dow Jones Indices, accessed April 27, 2026, https://www.spglobal.com/spdji/en/vix-intro/
65. A Regime-Switching Heston Model for VIX and S&P 500 Implied Volatilities - Princeton Economics Department, accessed April 27, 2026, https://economics.princeton.edu/published-papers/a-regime-switching-heston-model-for-vix-and-sp-500-implied-volatilities/
66. 1 Introduction - arXiv, accessed April 27, 2026, https://arxiv.org/html/2602.00080v1
67. Long Strangle Options Strategy: Beginner's Guide - TradingBlock, accessed April 27, 2026, https://www.tradingblock.com/strategies/long-strangle
68. Straddles/Strangles: Help me understand the math. : r/options - Reddit, accessed April 27, 2026, https://www.reddit.com/r/options/comments/1k8r31j/straddlesstrangles_help_me_understand_the_math/
69. Understanding Theta and Time Decay in Options | IBKR Quant, accessed April 27, 2026, https://www.interactivebrokers.com/campus/ibkr-quant-news/understanding-theta-and-time-decay-in-options/
70. Option Theta Explained: Time Decay for Beginners | TradingBlock, accessed April 27, 2026, https://www.tradingblock.com/blog/option-theta-time-decay
71. Long Straddle Options Strategy - Fidelity Investments, accessed April 27, 2026, https://www.fidelity.com/learning-center/investment-products/options/options-strategy-guide/long-straddle
72. Straddle vs Strangle Options: Which to Use and When? - PU Prime, accessed April 27, 2026, https://th.puprime.com/straddle-vs-strangle-options/
73. Using Straddles and Strangles to Manage Stock Events - Fidelity Investments, accessed April 27, 2026, https://www.fidelity.com/bin-public/060_www_fidelity_com/documents/learning-center/Deck_042821_Straddles-and-strangles.pdf
74. Straddles vs. Strangles Options Strategies - Charles Schwab, accessed April 27, 2026, https://www.schwab.com/learn/story/straddles-vs-strangles-options-strategies